



1.3.5: Statistical Tests and Models (In Person)

Session Objectives (*updated*)

1. Develop linear regression models and explore results.
 2. Develop logistic regression models and explore results.
 3. Perform t-tests and ANOVA and explore results.
 4. Modeling with Complex Survey Weights
-

0. Prework - Before You Begin

A. Install packages

If you do not have them already, install the following packages from CRAN (using the RStudio Menu “Tools/Install” Packages interface):

- [VIM](#) and [VIM package website](#)
- [gtsummary](#) and [gtsummary website](#)
- [easystats](#) and [easystats website](#)
- [car](#) and [car BOOK website](#)
- [effects](#)
- [olsrr](#) and [olsrr website](#)
- [dplyr](#) and [dplyr website](#)
- [ROCR](#) and [ROCR website](#)
- [effectsize](#) and [effectsize website](#)

**B. Open/create an RStudio project for this lesson**

Let's start with the `myfirstRproject` RStudio project you created in [Module 1.3.2 - part 1](#). If you have not yet created this `myfirstRproject` RStudio project, go ahead and create a new RStudio Project for this lesson. *Feel free to name your project whatever you want, it does not need to be named `myfirstRproject`.*



1. Develop linear regression models and explore results.

Linear Regression Modeling

As we saw briefly in [Module 1.3.4 - section 2 on missing data in regression models](#), linear regression can be accomplished using the built-in `lm()` function.

`lm()` stands for linear models. You can use this function for building both regression and ANOVA (analysis of variance) type models. `aov()` is another option for ANOVA as well.

Let's take a closer look at the little linear model we ran for the `sleep` dataset from the `VIM` package. We will run the regression model again and save the output to an object called `lm1`. Look at default output.

Using the `lm()` function and exploring output

```
# load VIM Package to get sleep dataset
library(VIM)

# run model for predicting "Sleep" from "Dream"
# save the output in lm1
lm1 <- lm(Sleep ~ Dream, data = sleep)

# look at default output
lm1
```

Call:

```
lm(formula = Sleep ~ Dream, data = sleep)
```

Coefficients:

(Intercept)	Dream
6.027	2.305



Let's take a moment to take a look at the `lm1` object in the Global Environment. Notice that only the intercept and slope terms are printed in the default output, but the `lm1` object is actually a list of 13 elements only one of which are the “coefficients” from the model.

Name	Type	Value
lm1	list [13] (S3: lm)	List of length 13
coefficients	double [2]	6.03 2.31
residuals	double [48]	-2.338 -6.277 2.159 4.683 -2.132 0.174 ...
effects	double [48]	-74.132 22.831 3.236 3.869 -1.233 -0.462 ...
rank	integer [1]	2
fitted.values	double [48]	10.64 10.18 7.64 15.02 8.33 14.33 ...
assign	integer [2]	0 1
qr	list [5] (S3: qr)	List of length 5
df.residual	integer [1]	46
na.action	integer [14] (S3: omit)	1 3 4 14 21 24 ...
xlevels	list [0]	List of length 0
call	language	lm(formula = Sleep ~ Dream, data = sleep)
terms	formula	Sleep ~ Dream
model	list [48 x 2] (S3: data.frame)	A data.frame with 48 rows and 2 columns



Next, to get more detailed output we need to run `summary(lm1)` which technically runs `summary.lm()` which is a special summary function specific for `lm` class type objects.

So, let's save the `summary(lm1)` output and also take a look at that object `slm1`.

```
# save the summary of lm1
slm1 <- summary(lm1)

# look at default output
slm1
```

Call:

```
lm(formula = Sleep ~ Dream, data = sleep)
```

Residuals:

Min	1Q	Median	3Q	Max
-6.2765	-2.0384	-0.1096	2.1599	9.2624

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	6.0273	0.7960	7.572	1.27e-09 ***
Dream	2.3051	0.3209	7.183	4.85e-09 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 3.178 on 46 degrees of freedom

(14 observations deleted due to missingness)

Multiple R-squared: 0.5287, Adjusted R-squared: 0.5184

F-statistic: 51.59 on 1 and 46 DF, p-value: 4.849e-09



But here is what is in the `slm1` object in the Global Environment - also a list with 12 elements. We get the `coefficients` again, but we get even more info - including the `df` (degrees of freedom), `r.squared` and `adj.r.squared`.

Name	Type	Value
slm1	list [12] (S3: summary.lm)	List of length 12
call	language	lm(formula = Sleep ~ Dream, data = sleep)
terms	formula	Sleep ~ Dream
residuals	double [48]	-2.338 -6.277 2.159 4.683 -2.132 0.174 ...
coefficients	double [2 x 4]	6.03e+00 2.31e+00 7.96e-01 3.21e-01 7.57e+00 7.18e+00 1.27e-09 4.85e-09 ...
aliased	logical [2]	FALSE FALSE
sigma	double [1]	3.178484
df	integer [3]	2 46 2
r.squared	double [1]	0.5286554
adj.r.squared	double [1]	0.5184088
fstatistic	double [3]	51.6 1.0 46.0
cov.unscaled	double [2 x 2]	0.0627 -0.0207 -0.0207 0.0102
na.action	integer [14] (S3: omit)	1 3 4 14 21 24 ...



Characteristic	Beta	95% CI ¹	p-value
Dream	2.3	1.7, 3.0	<0.001

¹CI = Confidence Interval

Nicer formatted regression table using `gtsummary::tbl_regression`

We can get a nicer output using the `gtsummary::tbl_regression()` table.

```
library(gtsummary)
tbl_regression(lm1) %>%
  as_gt() %>%
  tab_options(latex.tbl.pos = "h")
```



More output options for regression using easystats

Another suite of R packages that can be helpful to explore is the [easystats package suite](#). In this example we will look at:

- `model_parameters()` from the [parameters package](#) that is part of `easystats` package suite; and
- `report()` from the `[report package]` (<https://easystats.github.io/report/>) that is also part of the `easystats` package suite.

Better formatted output table:

```
library(easystats)
model_parameters(lm1)
```

Parameter	Coefficient	SE	95% CI	t(46)	p
(Intercept)	6.03	0.80	[4.43, 7.63]	7.57	< .001
Dream	2.31	0.32	[1.66, 2.95]	7.18	< .001



A nice summary with suggested interpretation verbiage:

```
report(lm1)
```

We fitted a linear model (estimated using OLS) to predict Sleep with Dream (formula: Sleep ~ Dream). The model explains a statistically significant and substantial proportion of variance ($R^2 = 0.53$, $F(1, 46) = 51.59$, $p < .001$, adj.

$R^2 = 0.52$). The model's intercept, corresponding to Dream = 0, is at 6.03 (95%

CI [4.43, 7.63], $t(46) = 7.57$, $p < .001$). Within this model:

- The effect of Dream is statistically significant and positive (beta = 2.31, 95% CI [1.66, 2.95], $t(46) = 7.18$, $p < .001$; Std. beta = 0.73, 95% CI [0.52, 0.93])

Standardized parameters were obtained by fitting the model on a standardized version of the dataset. 95% Confidence Intervals (CIs) and p-values were computed using a Wald t-distribution approximation.



Other options within tidyverse packages

In the [broom package](#) there are some additional functions that are helpful for exploring the model fit metrics and more.

```
library(broom)
tidy(lm1)
```

```
# A tibble: 2 x 5
  term          estimate std.error statistic    p.value
<chr>         <dbl>    <dbl>    <dbl>    <dbl>
1 (Intercept)    6.03    0.796     7.57 0.00000000127
2 Dream          2.31    0.321     7.18 0.00000000485
```

```
glance(lm1)
```

```
# A tibble: 1 x 12
  r.squared adj.r.squared sigma statistic    p.value    df logLik   AIC
  <dbl>      <dbl> <dbl>    <dbl>    <dbl> <dbl> <dbl> <dbl>
1 0.529      0.518  3.18     51.6 0.00000000485     1 -123. 251.
257.
# i 3 more variables: deviance <dbl>, df.residual <int>, nobs <int>
```


**“Companion for Applied Regression” - the car and effects packages**

John Fox has written an excellent set of books on applied regression that has an R companion book along with the `car` and `effects` packages with lots of helpful functions for doing regression modeling and analysis. Learn more at [Applied Regression Book](#) and [R Companion for Applied Regression Book](#).

Get the normal probability plot of the model residuals.

```
library(car)

# get normal probability plot of the
# regression model residuals
car::qqPlot(lm1$residuals)
```




33 42
24 32



Overlay a best fit line on the scatterplot of the original data for the model - include 95% confidence intervals for the best fit line.


```
# scatterplot of fitted model  
# using the car package, add the smooth option  
car::scatterplot(Sleep ~ Dream, data = sleep,  
                 smooth=TRUE)
```





Get an “effects” plot - for this model you only get one plot showing the slope of the line between Dream and Sleep:

```
library(effects)  
plot(allEffects(lm1))
```



**One more - the `olsrr` package**

The `olsrr` package provides a helpful set of tools for working with OLS (ordinary least squares) regression models. *Unfortunately, this set of package functions do NOT work with `glm` (generalized linear models) like logistic regression.*

Get detailed regression output including the standardized regression coefficients which are effect sizes, where `std.beta` = 0.1 is “small”, 0.3 is “moderate” and 0.5 is “large”.

```
# load olsrr package
library(olsrr)

# get detailed regression output
# including standardized coefficients
ols_regress(lm1)
```

Model Summary

R	0.727	RMSE	3.112
R-Squared	0.529	MSE	9.682
Adj. R-Squared	0.518	Coef. Var	29.705
Pred R-Squared	0.494	AIC	251.190
MAE	2.507	SBC	256.804

RMSE: Root Mean Square Error
MSE: Mean Square Error
MAE: Mean Absolute Error
AIC: Akaike Information Criteria
SBC: Schwarz Bayesian Criteria

ANOVA

	Sum of Squares	DF	Mean Square	F	Sig.
Regression	521.233	1	521.233	51.593	0.0000
Residual	464.727	46	10.103		
Total	985.960	47			

Parameter Estimates

model	Beta	Std. Error	Std. Beta	t	Sig.	lower	upper
-------	------	------------	-----------	---	------	-------	-------





(Intercept)	6.027	0.796		7.572	0.000	4.425
7.630						
Dream	2.305	0.321	0.727	7.183	0.000	1.659
2.951						

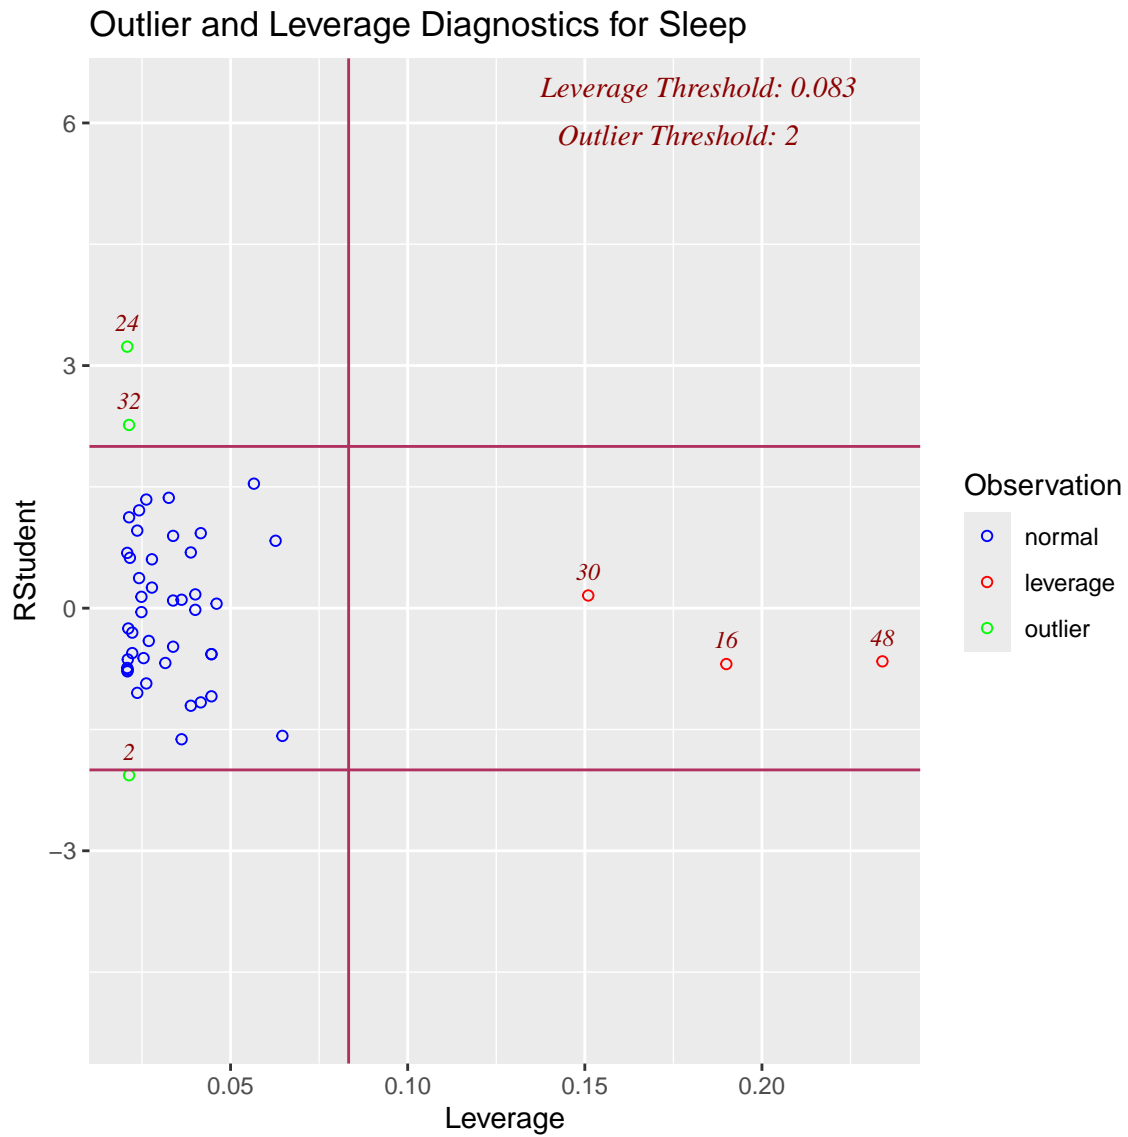



Get diagnostic plots.


```
# diagnostic plots  
# check for new output windows  
ols_plot_diagnostics(lm1)
```













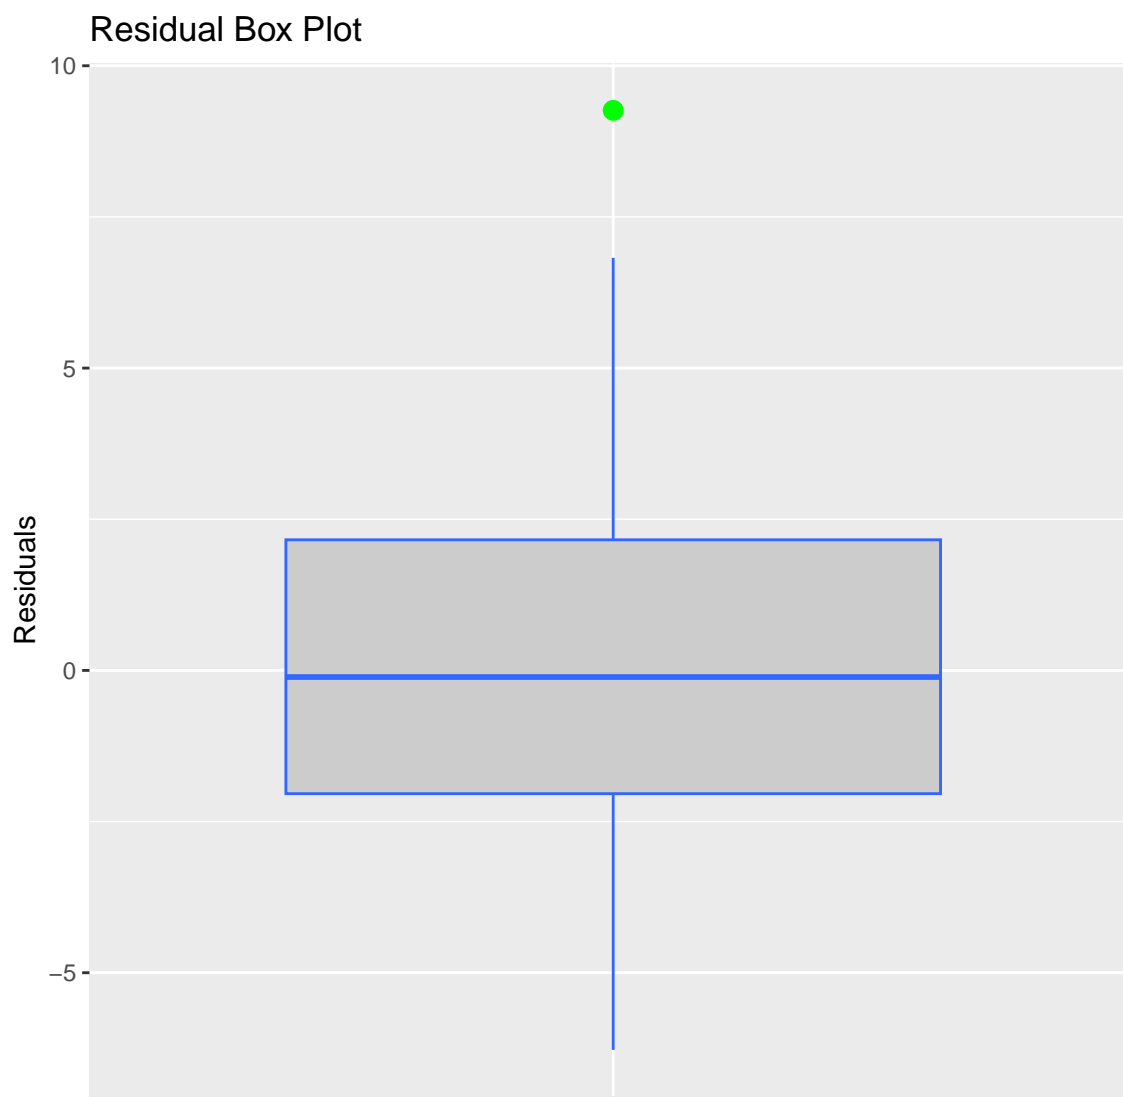






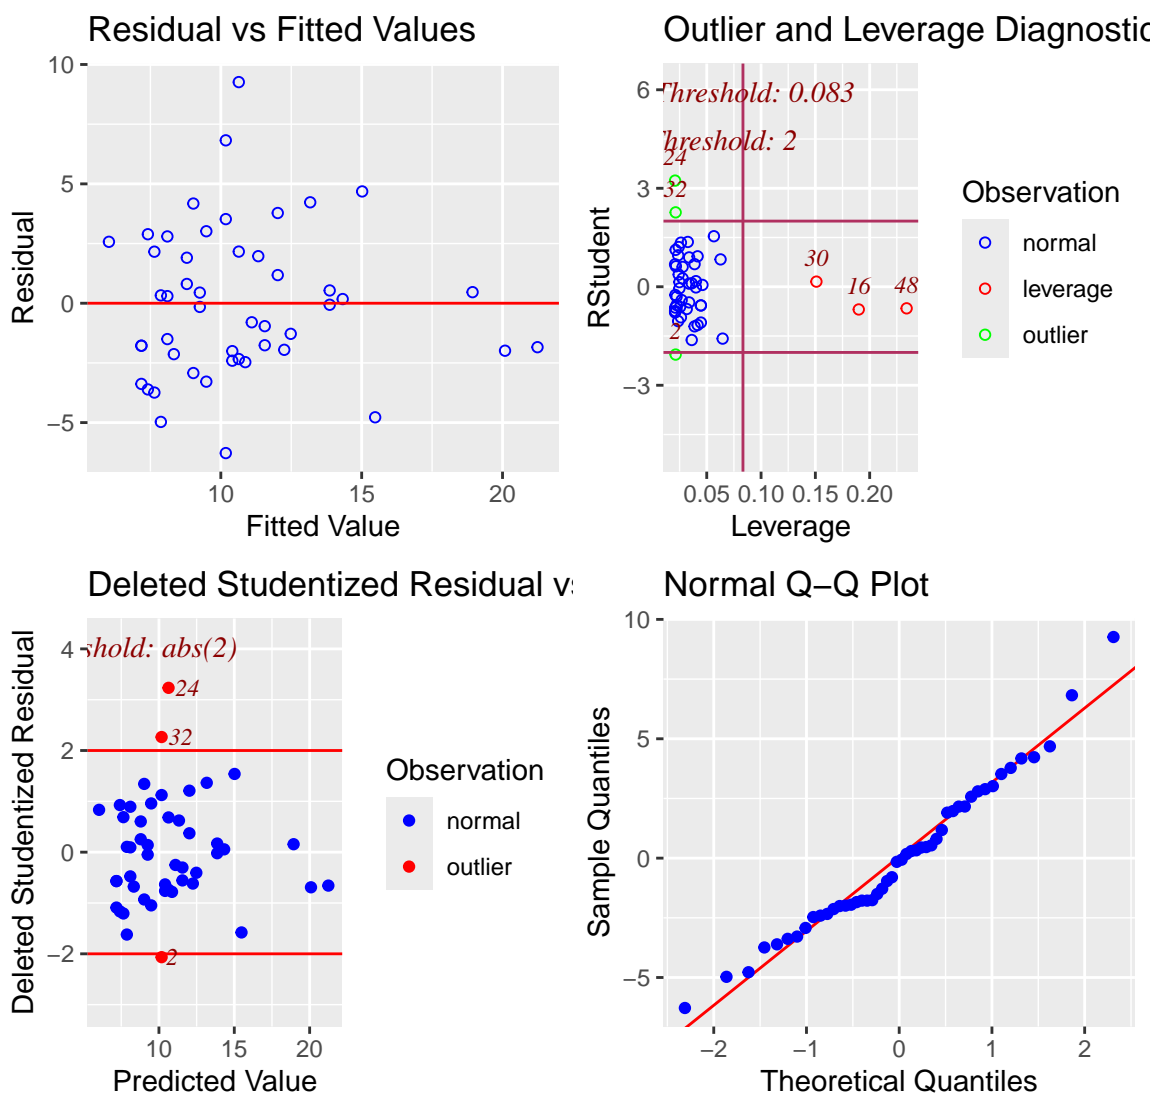









Regression Diagnostics






Regression Diagnostics

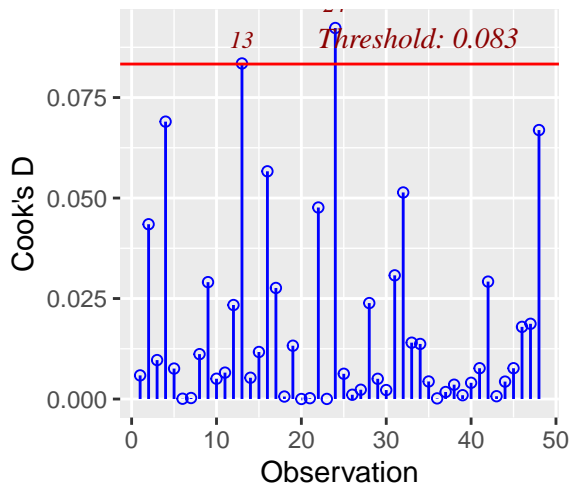
Actual vs Fitted for Sleep




Residual Fit Spread Plot



Cook's D Chart




Residual Fit Spread Plot






Regression Diagnostics

Residual Histogram



Residual Box Plot





Get a normality test for the residuals.

```
# normality tests for residuals  
ols_test_normality(lm1)
```

Test	Statistic	pvalue
Shapiro-Wilk	0.9736	0.3468
Kolmogorov-Smirnov	0.108	0.6307
Cramer-von Mises	3.1824	0.0000
Anderson-Darling	0.4289	0.2980



2. Develop logistic regression models and explore results.

Similar to what we did above, we can use `glm()` instead of `lm()` to perform a logistic regression. `glm()` is the generalized linear modeling function in base R. The “generalized” part of this is due to this function handling different “families” of output distributions. The “gaussian” family is the default for continuous variables with a normal distribution (the assumption for OLS). The family for logistic regression is the “binomial” since we are predicting the probability of someone being in group 1 or group 2 for the 2 possible outcomes in a logistic regression. And there are more families that can be fit including “poisson” which works for count-based variables (like number of children, miscarriages, etc).

This function can actually do a simple linear regression by leaving the default setting for `family = "gaussian"`. Learn more by running `help(glm, package = "stats")`.

Let’s keep working with the `sleep` dataset, but first let’s split the `Dream` variable into values ≤ 2 and those > 2 . *Note: The median for `Dream` was 1.8, so this should split data approximately 50/50.*

This time we will treat “`Dream > 2`” as the positive (or target) outcome for our logistic regression model. And then we will see how “`Dream > 2`” is predicted by amount of `Sleep` and `Danger` scores.

Simple `glm()` output

```
# create outcome variable
sleep$dream_gt2 <- as.numeric(sleep$Dream > 2)

# fit the logistic regression model
glm1 <- glm(dream_gt2 ~ Sleep + Danger,
            data = sleep,
            family = "binomial")

# look at basic output
glm1
```

```
Call:  glm(formula = dream_gt2 ~ Sleep + Danger, family = "binomial",
          data = sleep)
```

Coefficients:

(Intercept)	Sleep	Danger
-1.5953	0.2677	-0.8381



```
Degrees of Freedom: 47 Total (i.e. Null); 45 Residual  
(14 observations deleted due to missingness)  
Null Deviance:      63.51  
Residual Deviance: 41.79    AIC: 47.79
```

⚠ WARNING: Exponentiate Raw Coefficients

The default output from `glm()` only provides for the RAW logistic regression coefficients. To get the odds ratios, we must first exponentiate these coefficients.

Get the odds ratios by exponentiating the coefficients.

```
exp(coef(glm1))
```

(Intercept)	Sleep	Danger
0.2028428	1.3069544	0.4325214

**Get detailed glm() regression output.**

```
sglm1 <- summary(glm1)
sglm1
```

Call:

```
glm(formula = dream_gt2 ~ Sleep + Danger, family = "binomial",
     data = sleep)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.5953	1.5922	-1.002	0.3164
Sleep	0.2677	0.1151	2.325	0.0201 *
Danger	-0.8381	0.3888	-2.156	0.0311 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 63.510 on 47 degrees of freedom
Residual deviance: 41.795 on 45 degrees of freedom
(14 observations deleted due to missingness)
AIC: 47.795

Number of Fisher Scoring iterations: 5



Characteristic	OR ¹	95% CI ¹	p-value
Sleep	1.31	1.07, 1.70	0.020
Danger	0.43	0.18, 0.87	0.031

¹OR = Odds Ratio, CI = Confidence Interval

Use `gtsummary::tbl_regression()`

Get a nicer table and set `exponentiate = TRUE`.

```
tbl_regression(glm1, exponentiate = TRUE) %>%  
  as_gt() %>%  
  tab_options(latex.tbl.pos = "h")
```



Interpret Odds Ratios as Effect Sizes

Interpreting odds ratios as effect sizes is a little tricky. However, this website on [Computation of Effect Sizes](#) is really helpful - see items 14 and 16. By playing with this simple conversion tool to convert between effect sizes, we can see that:

- large effect sizes $d=0.8$, $r=0.5$, odds ratio $\approx 4.27-8.12$
- moderate effect sizes $d=0.5$, $r=0.3$, odds ratio $\approx 2.48-3.13$
- small effect sizes $d=0.2$, $r=0.1$, odds ratio ≈ 1.44

where Cohen's d is used for t-tests; r is used for correlations (and are the same for standardized regression coefficient "betas"); and odds ratios are from logistic regression.



AUC and ROC curve plot

The area under the curve (AUC) or “C-statistic” is often reported instead of R² for logistic regression models. The code below will compute the AUC for this model and make the receiver operating characteristic curve (ROC) plot.

Ideally you want AUC as close to 1 as possible:

- AUC > 0.9 is great
- AUC > 0.8 is good
- AUC > 0.7 is ok
- AUC < 0.7 is not very good
- AUC around 0.5 is no better than flipping a fair coin which is a useless model

As you can see below, the AUC for this model was 0.881 which is very good.

```
# NOTE: We need only the COMPLETE data
# for the 3 variables we used in this model
library(dplyr)
s1 <- sleep %>%
  select(Sleep, Danger, dream_gt2) %>%
  filter(complete.cases(.))


# compute AUC and get ROC curve
library(ROCR)
p <- predict(glm1, newdata=s1,
             type="response")
pr <- prediction(p, as.numeric(s1$dream_gt2))
prf <- performance(pr, measure = "tpr", x.measure = "fpr")

# compute AUC, area under the curve
# also called the C-statistic
auc <- performance(pr, measure = "auc")
auc <- auc@y.values[[1]]

# also - add title to plot with AUC in title
plot(prf,
     main = paste("ROC Curve, AUC = ", round(auc, 3)))
abline(0, 1, col="red")
```



ROC Curve, AUC = 0.881





3. Perform t-tests explore results.

T-tests

Given the split we did above looking at animals with `Dream` scores above and below 2, let's run a t-test for the `Sleep` variable.

```
# run t-test, save results
# default is an unequal variance "unpooled" t.test
tt1 <- t.test(Sleep ~ dream_gt2,
              data = sleep)
tt1
```

Welch Two Sample t-test

```
data: Sleep by dream_gt2
t = -4.5307, df = 38.191, p-value = 5.633e-05
alternative hypothesis: true difference in means between group 0 and group 1
is not equal to 0
95 percent confidence interval:
 -7.420174 -2.837604
sample estimates:
mean in group 0 mean in group 1
      8.776667      13.905556
```

```
# get the equal variance "pooled" t.test
tt2 <- t.test(Sleep ~ dream_gt2,
              data = sleep,
              var.equal = TRUE)
tt2
```

Two Sample t-test

```
data: Sleep by dream_gt2
t = -4.4417, df = 46, p-value = 5.566e-05
alternative hypothesis: true difference in means between group 0 and group 1
is not equal to 0
95 percent confidence interval:
 -7.453219 -2.804558
sample estimates:
```



```
mean in group 0 mean in group 1
      8.776667      13.905556
```



We can test the assumption of equal variance in a couple of ways. One simple way is to look at the standard deviations of each group and see if the ratio is larger than 2.

```
sd0 <- sleep %>%  
  filter(dream_gt2 == 0) %>%  
  select(Sleep) %>%  
  unlist() %>%  
  sd(na.rm = TRUE)  
sd1 <- sleep %>%  
  filter(dream_gt2 == 1) %>%  
  select(Sleep) %>%  
  unlist() %>%  
  sd(na.rm = TRUE)  
sd0
```

```
[1] 3.980615
```

```
sd1
```

```
[1] 3.682306
```

These standard deviations are similar, so a pooled t-test should be fine.



You can also run a formal test of equal variance, using `bartlett.test()`. However, this test and others like this are sensitive to small deviations from normality, so I often check the standard deviations and will run both the pooled and unpooled tests and see if I get the same conclusion either way.

As you can see, the p-value below is not significant, so we can not reject the null hypothesis assumption of equal variance - the pooled t-test is fine.

```
bartlett.test(Sleep ~ dream_gt2,  
              data = sleep)
```

Bartlett test of homogeneity of variances

data: Sleep by dream_gt2

Bartlett's K-squared = 0.12522, df = 1, p-value = 0.7234



Compute effect size for t-test

Use the [effectsize package](#) which is part of the `easystats` suite of packages.

The effect size computed is rather large, $d=1.32$.

```
library(effectsize)
options(es.use_symbols = TRUE)

cohens_d(Sleep ~ dream_gt2,
          data = sleep,
          na.action = na.omit)
```

Cohen's d	95% CI
-----------	--------

-1.32	[-1.96, -0.67]
-------	----------------

- Estimated using pooled SD.



Get a simple summary table

As you see below, the difference in Sleep scores between the 2 dream groups is approximately $13.9 - 8.8 = 5.1$ and the approximate average SD ≈ 3.9 , so the ratio of the mean differences to “pooled” SD $\approx 5.1/3.85 \approx 1.33$ which is close to the Cohen’s d we computed above.

I also added some custom statistical tests to the table.

```
# create factor variable with labels
sleep$dream_gt2.f <- factor(
  sleep$dream_gt2,
  levels = c(0, 1),
  labels = c("Dream <= 2",
             "Dream > 2")
)

tbl_summary(
  sleep,
  by = dream_gt2.f,
  include = c(Sleep, Danger),
  type = all_continuous() ~ "continuous2",
  statistic = all_continuous() ~ c("{N_nonmiss}", "{mean} ({sd})")
) %>%
  add_p(test = list(Sleep ~ "t.test", Danger ~ "wilcox.test"),
        test.args = list(Sleep ~ list(var.equal = TRUE))
  ) %>%
  as_gt() %>%
  tab_options(latex.tbl.pos = "h")
```



Characteristic	Dream ≤ 2 N = 32 ¹	Dream > 2 N = 18 ¹	p-value ²
Sleep			<0.001
N Non-missing	30	18	
Mean (SD)	8.8 (4.0)	13.9 (3.7)	
Unknown	2	0	
Danger			<0.001
1	5 (16%)	9 (50%)	
2	6 (19%)	6 (33%)	
3	5 (16%)	3 (17%)	
4	9 (28%)	0 (0%)	
5	7 (22%)	0 (0%)	

¹n (%)²Two Sample t-test; Wilcoxon rank sum test

4. Modeling with Complex Survey Weights

See [PRAMS Module](#)



R Code For This Module

- [module_135.R](#)

References

- Ben-Shachar, Mattan S., Daniel Lüdtke, and Dominique Makowski. 2020. “effectsize: Estimation of Effect Size Indices and Standardized Parameters.” *Journal of Open Source Software* 5 (56): 2815. <https://doi.org/10.21105/joss.02815>.
- Ben-Shachar, Mattan S., Dominique Makowski, Daniel Lüdtke, Indrajeet Patil, Brenton M. Wiernik, Rémi Thériault, and Philip Waggoner. 2025. *Effectsize: Indices of Effect Size*. <https://easystats.github.io/effectsize/>.
- Fox, John. 2003. “Effect Displays in R for Generalised Linear Models.” *Journal of Statistical Software* 8 (15): 1–27. <https://doi.org/10.18637/jss.v008.i15>.
- Fox, John, and Jangman Hong. 2009. “Effect Displays in R for Multinomial and Proportional-Odds Logit Models: Extensions to the effects Package.” *Journal of Statistical Software* 32 (1): 1–24. <https://doi.org/10.18637/jss.v032.i01>.
- Fox, John, and Sanford Weisberg. 2018. “Visualizing Fit and Lack of Fit in Complex Regression Models with Predictor Effect Plots and Partial Residuals.” *Journal of Statistical Software* 87 (9): 1–27. <https://doi.org/10.18637/jss.v087.i09>.
- . 2019a. *An R Companion to Applied Regression*. Third. Thousand Oaks CA: Sage. <https://www.john-fox.ca/Companion/>.
- . 2019b. *An R Companion to Applied Regression*. 3rd ed. Thousand Oaks CA: Sage. <https://socialsciences.mcmaster.ca/jfox/Books/Companion/index.html>.
- Fox, John, Sanford Weisberg, and Brad Price. 2022. *carData: Companion to Applied Regression Data Sets*. <https://r-forge.r-project.org/projects/car/>.
- . 2024. *Car: Companion to Applied Regression*. <https://r-forge.r-project.org/projects/car/>.
- Fox, John, Sanford Weisberg, Brad Price, Michael Friendly, and Jangman Hong. 2022. *Effects: Effect Displays for Linear, Generalized Linear, and Other Models*. <https://www.r-project.org>.
- Hebbali, Aravind. 2024. *Olsrr: Tools for Building OLS Regression Models*. <https://olsrr.rsquaredacademy.com/>.
- Iannone, Richard, Joe Cheng, Barret Schloerke, Ellis Hughes, Alexandra Lauer, JooYoung Seo, Ken Brevoort, and Olivier Roy. 2024. *Gt: Easily Create Presentation-Ready Display Tables*. <https://gt.rstudio.com>.
- Ihaka, Ross, Paul Murrell, Kurt Hornik, Jason C. Fisher, Reto Stauffer, Claus O. Wilke, Claire D. McWhite, and Achim Zeileis. 2023. *Colorspace: A Toolbox for Manipulating and Assessing Colors and Palettes*. <https://colorspace.R-Forge.R-project.org/>.
- Kowarik, Alexander, and Matthias Templ. 2016. “Imputation with the R Package VIM.” *Journal of Statistical Software* 74 (7): 1–16. <https://doi.org/10.18637/jss.v074.i07>.



- Lüdecke, Daniel, Mattan S. Ben-Shachar, Indrajeet Patil, and Dominique Makowski. 2020. “Extracting, Computing and Exploring the Parameters of Statistical Models Using R.” *Journal of Open Source Software* 5 (53): 2445. <https://doi.org/10.21105/joss.02445>.
- Lüdecke, Daniel, Mattan S. Ben-Shachar, Indrajeet Patil, Philip Waggoner, and Dominique Makowski. 2021. “performance: An R Package for Assessment, Comparison and Testing of Statistical Models.” *Journal of Open Source Software* 6 (60): 3139. <https://doi.org/10.21105/joss.03139>.
- Lüdecke, Daniel, Mattan S. Ben-Shachar, Indrajeet Patil, Brenton M. Wiernik, Etienne Bacher, Rémi Thériault, and Dominique Makowski. 2022. “Easystats: Framework for Easy Statistical Modeling, Visualization, and Reporting.” *CRAN*. <https://doi.org/10.32614/CRAN.package.easystats>.
- Lüdecke, Daniel, Dominique Makowski, Mattan S. Ben-Shachar, Indrajeet Patil, Søren Højsgaard, and Brenton M. Wiernik. 2025. *Parameters: Processing of Model Parameters*. <https://easystats.github.io/parameters/>.
- Lüdecke, Daniel, Dominique Makowski, Mattan S. Ben-Shachar, Indrajeet Patil, Philip Waggoner, Brenton M. Wiernik, and Rémi Thériault. 2025. *Performance: Assessment of Regression Models Performance*. <https://easystats.github.io/performance/>.
- Lüdecke, Daniel, Dominique Makowski, Mattan S. Ben-Shachar, Indrajeet Patil, Brenton M. Wiernik, Etienne Bacher, and Rémi Thériault. 2025. *Easystats: Framework for Easy Statistical Modeling, Visualization, and Reporting*. <https://easystats.github.io/easystats/>.
- Lüdecke, Daniel, Dominique Makowski, Indrajeet Patil, Mattan S. Ben-Shachar, Brenton M. Wiernik, and Philip Waggoner. 2025. *See: Model Visualisation Toolbox for Easystats and Ggplot2*. <https://easystats.github.io/see/>.
- Lüdecke, Daniel, Dominique Makowski, Indrajeet Patil, Philip Waggoner, Mattan S. Ben-Shachar, Brenton M. Wiernik, Vincent Arel-Bundock, and Etienne Bacher. 2025. *Insight: Easy Access to Model Information for Various Model Objects*. <https://easystats.github.io/insight/>.
- Lüdecke, Daniel, Indrajeet Patil, Mattan S. Ben-Shachar, Brenton M. Wiernik, Philip Waggoner, and Dominique Makowski. 2021. “see: An R Package for Visualizing Statistical Models.” *Journal of Open Source Software* 6 (64): 3393. <https://doi.org/10.21105/joss.03393>.
- Lüdecke, Daniel, Philip Waggoner, and Dominique Makowski. 2019. “insight: A Unified Interface to Access Information from Model Objects in R.” *Journal of Open Source Software* 4 (38): 1412. <https://doi.org/10.21105/joss.01412>.
- Makowski, Dominique, Mattan S. Ben-Shachar, and Daniel Lüdecke. 2019. “bayestestR: Describing Effects and Their Uncertainty, Existence and Significance Within the Bayesian Framework.” *Journal of Open Source Software* 4 (40): 1541. <https://doi.org/10.21105/joss.01541>.
- Makowski, Dominique, Mattan S. Ben-Shachar, Indrajeet Patil, and Daniel Lüdecke. 2020. “Methods and Algorithms for Correlation Analysis in R.” *Journal of Open Source Software* 5 (51): 2306. <https://doi.org/10.21105/joss.02306>.
- Makowski, Dominique, Mattan S. Ben-Shachar, Brenton M. Wiernik, Indrajeet Patil, Rémi Thériault, and Daniel Lüdecke. 2025. “modelbased: An R Package to Make the Most



- Out of Your Statistical Models Through Marginal Means, Marginal Effects, and Model Predictions.” *Journal of Open Source Software* 10 (109): 7969. <https://doi.org/10.21105/joss.07969>.
- Makowski, Dominique, Daniel Lüdtke, Mattan S. Ben-Shachar, Indrajeet Patil, and Rémi Thériault. 2025. *Modelbased: Estimation of Model-Based Predictions, Contrasts and Means*. <https://easystats.github.io/modelbased/>.
- Makowski, Dominique, Daniel Lüdtke, Mattan S. Ben-Shachar, Indrajeet Patil, Micah K. Wilson, and Brenton M. Wiernik. 2025. *bayestestR: Understand and Describe Bayesian Models and Posterior Distributions*. <https://easystats.github.io/bayestestR/>.
- Makowski, Dominique, Daniel Lüdtke, Indrajeet Patil, Rémi Thériault, Mattan S. Ben-Shachar, and Brenton M. Wiernik. 2023. “Automated Results Reporting as a Practical Tool to Improve Reproducibility and Methodological Best Practices Adoption.” *CRAN*. <https://easystats.github.io/report/>.
- . 2025. *Report: Automated Reporting of Results and Statistical Models*. <https://easystats.github.io/report/>.
- Makowski, Dominique, Brenton M. Wiernik, Indrajeet Patil, Daniel Lüdtke, and Mattan S. Ben-Shachar. 2022. “correlation: Methods for Correlation Analysis.” <https://CRAN.R-project.org/package=correlation>.
- Makowski, Dominique, Brenton M. Wiernik, Indrajeet Patil, Daniel Lüdtke, Mattan S. Ben-Shachar, and Rémi Thériault. 2025. *Correlation: Methods for Correlation Analysis*. <https://easystats.github.io/correlation/>.
- Patil, Indrajeet, Etienne Bacher, Dominique Makowski, Daniel Lüdtke, Mattan S. Ben-Shachar, and Brenton M. Wiernik. 2025. *Datawizard: Easy Data Wrangling and Statistical Transformations*. <https://easystats.github.io/datawizard/>.
- Patil, Indrajeet, Dominique Makowski, Mattan S. Ben-Shachar, Brenton M. Wiernik, Etienne Bacher, and Daniel Lüdtke. 2022. “datawizard: An R Package for Easy Data Preparation and Statistical Transformations.” *Journal of Open Source Software* 7 (78): 4684. <https://doi.org/10.21105/joss.04684>.
- R Core Team. 2025. *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing. <https://www.R-project.org/>.
- Robinson, David, Alex Hayes, and Simon Couch. 2025. *Broom: Convert Statistical Objects into Tidy Tibbles*. <https://broom.tidymodels.org/>.
- Sing, Tobias, Oliver Sander, Niko Beerenwinkel, and Thomas Lengauer. 2020. *ROCR: Visualizing the Performance of Scoring Classifiers*. <http://ipa-tys.github.io/ROCR/>.
- Sing, T., O. Sander, N. Beerenwinkel, and T. Lengauer. 2005. “ROCR: Visualizing Classifier Performance in r.” *Bioinformatics* 21 (20): 7881. <http://rocr.bioinf.mpi-sb.mpg.de>.
- Sjoberg, Daniel D., Joseph Larmarange, Michael Curry, Jessica Lavery, Karissa Whiting, and Emily C. Zabor. 2024. *Gtsummary: Presentation-Ready Data Summary and Analytic Result Tables*. <https://github.com/ddsjoberg/gtsummary>.
- Sjoberg, Daniel D., Karissa Whiting, Michael Curry, Jessica A. Lavery, and Joseph Larmarange. 2021. “Reproducible Summary Tables with the Gtsummary Package.” *The R Journal* 13: 570–80. <https://doi.org/10.32614/RJ-2021-053>.
- Stauffer, Reto, Georg J. Mayr, Markus Dabernig, and Achim Zeileis. 2009. “Somewhere



over the Rainbow: How to Make Effective Use of Colors in Meteorological Visualizations.” *Bulletin of the American Meteorological Society* 96 (2): 203–16. <https://doi.org/10.1175/BAMS-D-13-00155.1>.

Templ, Matthias, Alexander Kowarik, Andreas Alfons, Gregor de Cillia, and Wolfgang Rannetbauer. 2022. *VIM: Visualization and Imputation of Missing Values*. <https://github.com/statistikat/VIM>.

Wickham, Hadley, Romain François, Lionel Henry, Kirill Müller, and Davis Vaughan. 2023. *Dplyr: A Grammar of Data Manipulation*. <https://dplyr.tidyverse.org>.

Zeileis, Achim, Jason C. Fisher, Kurt Hornik, Ross Ihaka, Claire D. McWhite, Paul Murrell, Reto Stauffer, and Claus O. Wilke. 2020. “colorspace: A Toolbox for Manipulating and Assessing Colors and Palettes.” *Journal of Statistical Software* 96 (1): 1–49. <https://doi.org/10.18637/jss.v096.i01>.

Zeileis, Achim, Kurt Hornik, and Paul Murrell. 2009. “Escaping RGBland: Selecting Colors for Statistical Graphics.” *Computational Statistics & Data Analysis* 53 (9): 3259–70. <https://doi.org/10.1016/j.csda.2008.11.033>.

Other Helpful Resources

[Other Helpful Resources](#)